

AgentOps Supply Chain Governance Policy Handbook

Enterprise Policy Document

Effective Date: June 1, 2026

Version: 1.0

1. Document Control

Document Title: AgentOps Supply Chain Governance Policy Handbook

Version: 1.0

Effective Date: June 1, 2026

Last Updated: June 1, 2026

Classification: Internal Use

Approval Authority: Supply Chain Leadership

2. Purpose and Scope

This policy handbook provides governance guidance for the AgentOps Control Tower, an autonomous supply chain management system that uses multiple AI agents to make operational decisions. The handbook establishes clear policies, decision criteria, risk thresholds, and approval workflows to ensure that autonomous agent actions remain within acceptable risk parameters and organizational standards.

This handbook applies to all autonomous agents operating within the supply chain domain, including forecasting, inventory, procurement, and logistics agents. It governs agent access to datasets, tool usage, financial decisions, external communications, and escalation workflows.

3. Definitions

Autonomous Agent: An AI system that makes recommendations and decisions within defined governance constraints.

Policy Decision: A determination of whether an agent action should be Allowed, Escalated, or Blocked.

Escalation: A request for human review and approval before action execution.

Block: A prohibition on action execution due to policy violation or risk threshold breach.

Risk Score: A numeric assessment (0-100) of decision risk based on policy factors.

Restricted Data: Sensitive datasets (HR, payroll, PII) that agents may not access.

Source Traceability: Documentation of data sources and record IDs used in a decision.

4. Agent Roles and Responsibilities

The AgentOps system includes multiple specialized agents, each with defined responsibilities and constraints:

Coordinator Agent: Orchestrates workflows and routes decisions to appropriate agents. Operates at orchestration level with broad dataset access and escalation authority.

Forecasting Agent: Generates demand forecasts based on sales history. Restricted to products and sales history datasets. Provides confidence scores with all forecasts.

Inventory Agent: Monitors stock levels and calculates shortage risk. Accesses product and inventory data. Escalates critical shortages.

Procurement Agent: Recommends suppliers and purchase quantities. Subject to high-value approval threshold and supplier compliance validation.

Logistics Agent: Optimizes routes and evaluates disruption impacts. May escalate for high-disruption route selections.

Policy Engine: Enforces governance rules. Validates all agent actions against this handbook.

Risk Scoring Engine: Calculates risk scores based on policy factors. Recommends escalation if risk exceeds thresholds.

5. Data Access Governance Policy

Each agent is assigned an explicit list of datasets it may access for decision-making. Agents must not exceed these boundaries. Dataset access is enforced at the point of agent invocation and during policy evaluation.

Agents shall only access datasets necessary for their assigned function. The Coordinator Agent has access to the broadest dataset set due to its orchestration role. Specialized agents (Forecasting, Inventory, Procurement, Logistics) have restricted access aligned with their domain.

Quarterly reviews of agent dataset access are conducted to ensure continued appropriateness. Any changes to dataset access require explicit approval from the Policy Authority.

6. Tool Usage Governance Policy

Similar to dataset access, each agent has an assigned list of tools it is authorized to use. These include system functions such as supplier selection, route optimization, forecast generation, and approval routing.

Agents must not invoke unauthorized tools. Any attempt to use a tool outside the assigned toolset is blocked. New tool authorization requires explicit policy amendment and audit trail documentation.

7. Procurement Governance Policy

Procurement decisions are subject to value-based approval thresholds and supplier compliance validation. The goal is to balance operational autonomy with financial and vendor risk management.

Financial Authority: Autonomous agents may execute low-value procurement (below INR 50,000) without human approval, provided all other policy conditions are met. Procurement recommendations above this threshold must be escalated for human review before execution.

Execution Restriction: Agents may recommend procurement but must not finalize, execute, or externally communicate the purchase without appropriate approval.

8. Supplier Compliance Policy

Supplier selection is governed by compliance and approval status. The organization maintains a supplier database with approval and compliance attributes for each vendor.

Approved and Compliant: Suppliers with approval status "Yes" and compliance status "Compliant" are preferred choices for procurement. These suppliers have passed due diligence and audit reviews.

Approved but Under Review: Suppliers with status "Under Review" may be considered only when no compliant alternatives are available. Selection of such suppliers may trigger additional review depending on the overall risk score.

Non-Compliant or Unapproved: Suppliers marked as "Non-Compliant" or with approval status "No" must not be selected. If an agent recommends a non-compliant or unapproved supplier, the action must be blocked.

9. Logistics and Route Disruption Policy

Logistics decisions include route selection, transportation mode choice, and disruption impact assessment. Routes may have active disruptions (weather, congestion, infrastructure issues) that affect delivery time, cost, and risk.

Active Disruption Assessment: Before finalizing a route selection, the Logistics Agent must check for active disruptions on the selected route. Active disruptions are those with status "Active" and have start and end dates that include the current date or planned delivery window.

High or Critical Severity: If a selected route has an active disruption with severity rated "High" or "Critical", the recommendation must be escalated for human review. High-severity disruptions may significantly increase delivery delays, operational costs, or business continuity risk.

Route Alternatives: The Logistics Agent should maintain a ranked list of alternative routes. If the preferred route has an active high-severity disruption, the agent should recommend alternatives and escalate for decision.

10. Source Traceability Policy

Every agent recommendation must include source traceability. This means identifying the data sources, specific records, and policy references that informed the decision. Source traceability provides auditability, explainability, and enables verification of decision logic.

Required Documentation: Final recommendations must cite source files (e.g., sales_history.csv, suppliers.csv), source record identifiers (e.g., product_id, supplier_id, date range), and relevant policy clauses.

Missing Traceability: If a final recommendation lacks source documentation, it must be blocked. This ensures that all decisions can be audited and explained.

RAG Agent Evidence: When the Policy RAG Agent retrieves policy text from this handbook, it must cite the specific policy sections retrieved and include them in the decision evidence.

11. External Communication Policy

Autonomous agents must not initiate external communications to suppliers, vendors, customers, or external parties without human approval. External communications include supplier emails, purchase order submissions, vendor API calls, and external notifications.

MVP Constraint: During the MVP phase, all external communications are prohibited from autonomous agents. Any agent action that includes or requires external communication must be escalated for human execution.

Future State: Post-MVP, external communication may be authorized for selected agents and communication types, subject to approval routing and audit logging.

12. Agent Status Governance Policy

Each agent has an operational status: Active, Inactive, or Suspended. Only Active agents may execute decisions, access tools, or make recommendations.

Active Status: The agent is authorized to operate and may execute decisions within policy constraints.

Inactive Status: The agent is temporarily unavailable. Any action from an inactive agent must be blocked.

Suspended Status: The agent is prohibited from operating, typically due to policy violations or security concerns. Any action from a suspended agent must be blocked.

13. Risk Scoring and Human Approval Policy

All agent actions are evaluated for risk using a scoring framework. Risk scores range from 0 to 100, with higher scores indicating higher risk. Risk scores are calculated based on policy factors such as procurement value, supplier approval, data access restrictions, and route disruptions.

Risk Thresholds:

Low Risk (0-30): No escalation required; action may proceed if no blocking policies apply.

Medium Risk (31-60): Action may proceed if no blocking policies apply; monitoring recommended.

High Risk (61-80): Action must be escalated for human review if no blocking policies apply.

Critical Risk (81-100): Action must be escalated for human review and requires explicit approval.

Approval Queue: Actions classified as Escalate are placed into a human approval queue. A designated reviewer (e.g., Supply Chain Manager) must review, approve, or reject the action. The action must not be executed until approval is granted.

Approval Timeout: Escalated actions should be reviewed within four business hours. If no approval is received within this timeframe, the escalation should trigger a secondary review.

14. Audit Logging and Evidence Retention Policy

Complete audit trails are maintained for all agent actions and policy decisions. Audit logs enable compliance verification, incident investigation, and system improvement.

Required Log Elements: Each audit log entry must include timestamp, agent ID, action type, datasets accessed, tools used, policy conditions evaluated, risk score, approval status, source references, final decision (Allow/Escalate/Block), and policy evidence retrieved.

Retention Period: Audit logs are retained for a minimum of two years. Logs related to escalated or blocked actions are retained for three years.

Access Control: Audit logs are accessible to authorized compliance and security personnel only.

15. Enforcement Actions

Policy violations result in specific enforcement actions determined by the severity of the violation:

First Violation: Alert logged; agent continues operation under heightened monitoring; policy refresh conducted.

Second Violation: Agent status changed to Inactive; incident reviewed by policy authority; requires explicit reactivation approval.

Repeated or Critical Violation: Agent status changed to Suspended; full audit of agent actions conducted; escalated to leadership for investigation.

16. Policy Interpretation Rules

Decision Priority: When multiple policies apply to a single action, the final decision follows this priority order: Block overrides Escalate. Escalate overrides Allow. Allow applies only when no blocking or escalation rules are triggered.

Condition Matching: Policy conditions must match actual system state exactly. Partial or ambiguous matches should result in escalation rather than automatic allow.

Conservative Interpretation: In cases of policy ambiguity or low-confidence interpretation, the system defaults to escalation for human review.

Evidence Requirement: If the Policy RAG Agent cannot retrieve relevant policy evidence with sufficient confidence, the action must be escalated rather than automatically allowed.

17. Example Policy Scenarios

The following scenarios demonstrate how policy rules are applied in practice:

Scenario 1: High-Value Procurement Request

Situation: The Procurement Agent recommends purchasing a critical component at a total value of INR 387,000 from an approved supplier with no disruptions on the route.

Policy Analysis: The procurement value (INR 387,000) exceeds the authorization threshold of INR 50,000. Policy 7 (Procurement Governance) applies.

Decision: ESCALATE. The action must be placed in the human approval queue. The Supply Chain Manager reviews the recommendation, including supplier reputation, inventory urgency, and business justification. Upon approval, the procurement proceeds. If rejected, the agent is notified and alternative options are considered.

Scenario 2: Unapproved Supplier Selection

Situation: The Procurement Agent recommends selecting a supplier that has approval status "No" and compliance status "Non-Compliant".

Policy Analysis: Policy 8 (Supplier Compliance) applies. The selected supplier is marked as unapproved.

Decision: BLOCK. The recommendation is rejected. The system logs the policy violation and notifies the Procurement Agent. The agent must select an approved alternative supplier. If no approved suppliers exist for the required component, the situation must be escalated for exception handling.

Scenario 3: Active Route Disruption

Situation: The Logistics Agent recommends a road route for an urgent supplier shipment. The disruption database shows this route has an active disruption with severity "High" and status "Active".

Policy Analysis: Policy 9 (Logistics and Route Disruption) applies. The route has an active High-severity disruption.

Decision: ESCALATE. The recommendation is escalated to a human reviewer (Logistics Manager). The reviewer considers alternative routes, cost trade-offs, and delivery deadline impact. If the disruption is short-duration and alternatives are insufficient, the manager may approve the original route despite the disruption.

Scenario 4: Restricted Data Access Attempt

Situation: The Forecasting Agent is processing a demand forecast. During data access validation, the system detects that the agent attempted to query the payroll.csv file.

Policy Analysis: Policy 5 (Data Access Governance) applies. The Forecasting Agent is not authorized to access payroll.csv.

Decision: BLOCK. The data access request is denied. The action is logged as a policy violation. The Forecasting Agent is notified that the forecast cannot be completed using restricted data.

Scenario 5: Missing Source Traceability

Situation: The Coordinator Agent returns a final procurement recommendation. The recommendation text states "Procure 1000 units from Supplier S-012." However, the recommendation does not include any source files, source record IDs, or justification data.

Policy Analysis: Policy 10 (Source Traceability) applies. The recommendation lacks source documentation.

Decision: BLOCK. The recommendation is rejected due to missing traceability. The system notifies the Coordinator Agent to resubmit the recommendation with complete source citations, including the inventory records used, sales history data, and supplier selection criteria.

Scenario 6: External Communication Attempt

Situation: The Procurement Agent has approved a supplier recommendation and generates a draft purchase order email to send to the supplier. The system detects this external communication attempt.

Policy Analysis: Policy 11 (External Communication) applies. The agent attempted to send external communication.

Decision: ESCALATE. External communication is escalated to a human (Procurement Manager) for execution. The manager reviews the draft email for appropriateness, accuracy, and legal compliance before sending it.

18. Appendix: Policy-to-System Mapping

This appendix defines the relationship between policy clauses and system implementation. It specifies the expected condition fields, actions, and severity levels that the RAG-based Policy Engine extracts from this handbook.

Policy Area: High-Value Procurement Approval

Condition Field: procurement_value

Condition: procurement_value is greater than INR 50,000

Expected Action: Escalate

Severity: High

Evidence Requirement: Cite Section 7 (Procurement Governance Policy) and the clause defining the INR 50,000 threshold.

Policy Area: Unapproved Supplier Block

Condition Field: is_approved

Condition: is_approved equals "No" OR compliance_status equals "Non-Compliant"

Expected Action: Block

Severity: Critical

Evidence Requirement: Cite Section 8 (Supplier Compliance Policy). Include the supplier approval status and compliance status.

Policy Area: Supplier Compliance Preference

Condition Field: compliance_status

Condition: compliance_status equals "Compliant" (preferred) OR "Under Review" (conditional)

Expected Action: Allow (if compliant); Escalate or conditional allow (if under review)

Severity: Medium (if under review)

Evidence Requirement: Cite Section 8 (Supplier Compliance Policy) with explanation of compliance status implications.

Policy Area: Route Disruption Impact

Condition Field: route_disruption_status, route_disruption_severity

Condition: route_disruption_status equals "Active" AND (route_disruption_severity equals "High" OR "Critical")

Expected Action: Escalate

Severity: High

Evidence Requirement: Cite Section 9 (Logistics and Route Disruption Policy). Include disruption ID, route ID, disruption type, severity, and expected impact.

Policy Area: Restricted Data Access Block

Condition Field: restricted_data_accessed

Condition: restricted_data_accessed equals true (e.g., hr_data.csv, payroll.csv, employee_records.csv, customer_pii.csv)

Expected Action: Block

Severity: Critical

Evidence Requirement: Cite Section 5 (Data Access Governance Policy). Specify which restricted dataset was accessed.

Policy Area: Source Traceability Requirement

Condition Field: source_citation_missing

Condition: source_citation_missing equals true (no source files, record IDs, or policy evidence included)

Expected Action: Block

Severity: High

Evidence Requirement: Cite Section 10 (Source Traceability Policy). Explain that all recommendations must include source documentation.

Policy Area: External Communication Escalation

Condition Field: external_communication_attempted

Condition: external_communication_attempted equals true

Expected Action: Escalate

Severity: High

Evidence Requirement: Cite Section 11 (External Communication Policy). Include communication type and intended recipient.

Policy Area: Unauthorized Tool Use

Condition Field: unauthorized_tool_used

Condition: unauthorized_tool_used equals true (agent used a tool not in its assigned toolset)

Expected Action: Block

Severity: High

Evidence Requirement: Cite Section 6 (Tool Usage Governance Policy). Reference the agent's assigned toolset.

Policy Area: Agent Status Governance

Condition Field: agent_status

Condition: agent_status not equals "Active" (i.e., "Inactive" or "Suspended")

Expected Action: Block

Severity: Critical

Evidence Requirement: Cite Section 12 (Agent Status Governance Policy). Include agent ID

and current status.

Policy Area: Forecast Confidence Risk

Condition Field: forecast_confidence

Condition: forecast_confidence is less than 0.70

Risk Factor: Low confidence contributes 20 points to risk score; does not block automatically but increases risk level

Expected Action: Allow (if risk does not trigger escalation); Escalate (if total risk is High or Critical)

Severity: Medium (confidence factor only)

Evidence Requirement: Cite Section 13 (Risk Scoring and Human Approval Policy). Include confidence score and note that low confidence increases overall risk.

Policy Area: Decision Priority Logic

Decision Rule: When multiple policies apply, follow this priority:

1. If any policy triggers "Block", the final decision is BLOCK.
2. Else if any policy triggers "Escalate", the final decision is ESCALATE.
3. Else if risk-based escalation is enabled and risk level is High or Critical, the final decision is ESCALATE.
4. Else the final decision is ALLOW.

Evidence Requirement: Cite Section 16 (Policy Interpretation Rules). Explain which policies were evaluated.

Policy Area: Evidence Requirement and Conservative Interpretation

Condition Field: policy_evidence_confidence

Condition: If the Policy RAG Agent cannot retrieve relevant policy evidence with sufficient confidence, or if policy interpretation is ambiguous

Expected Action: Escalate (not Allow)

Severity: Medium to High (depending on decision impact)

Evidence Requirement: Cite Section 16 (Policy Interpretation Rules). Explain that conservative interpretation defaults to escalation when policy evidence is unclear.